

Facial Expression Recognition

Emmeline Tsen

Jacky Chow

Jerry Huang

Abstract

- Human face conveys a lot of information
- Facial detection has been a popular topic within AI
- Determine a person's facial expression in an image
 - Help better understand what the person is feeling based on the image



Introduction

- Facial Expression Recognition (FER) has received a lot of attention in recent years within several fields
 - Medical
 - Security
 - Communications
 - Etc.
- Project consists of:
 - CNN Model & Model Tuning
 - Generative Adversarial Networks (GANs) to generate faces
 - CNN model to determine emotions from faces generated using GAN
 - TFX for Tensorflow Serving Model



Related Work

- Mainstream methods
 - Traditional manual methods
 - Local Binary Patterns (LBP)
 - Shallow learning
 - SVM
 - Adaboost
 - Deep learning
 - CNN
 - RNN



Data

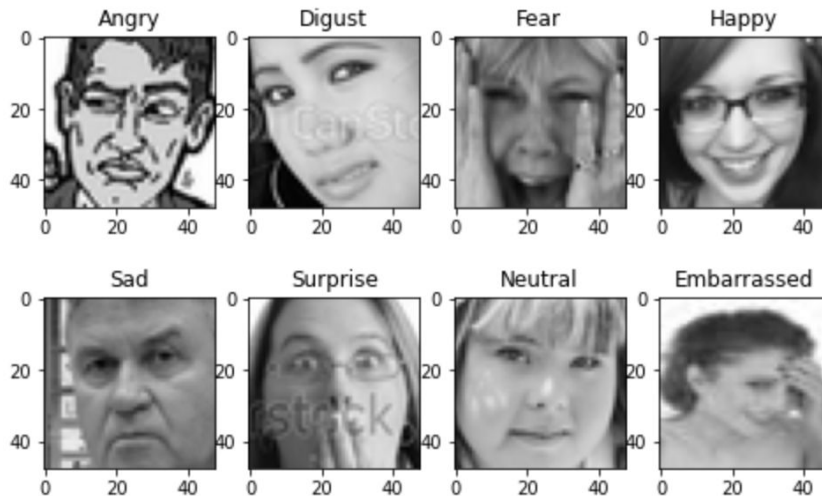
- Dataset

- Fer2013 dataset
- Amalgamation
 - Google image search:
Embarrassed males and females

- Data Pre-processing

- Image to Array
- Converted images to black and white
- Converted images to $48 * 48$

```
show_img(data_valid)
```



Methods -- CNN

- Reshape the data into the fitting size (4 dimension array)
- Build the model
 - With data generator (rotation, shift, zoom)
 - Early stopping
- Hyperparameter Tuning
 - BatchNormalization
 - Dropout
 - Learning Rate
 - CNN Layers
 - Epochs
 - Optimizer

Colab - https://github.com/emmelinetsen/faceAI/blob/master/Main_CNN_FER.ipynb




Experiments

CNN Model	Drop out	Zoom	Epochs	Batch Normalize	Max Accuracy
1	0.25	yes	25	no	0.60
2	0.25	no	50	yes	0.61
3	0.25 (each layer)	no	50	yes	0.60
4	0.5	no	75	yes	0.64
5	0.75	no	50	yes	0.25

Best Model

Accuracy

epoch_accuracy

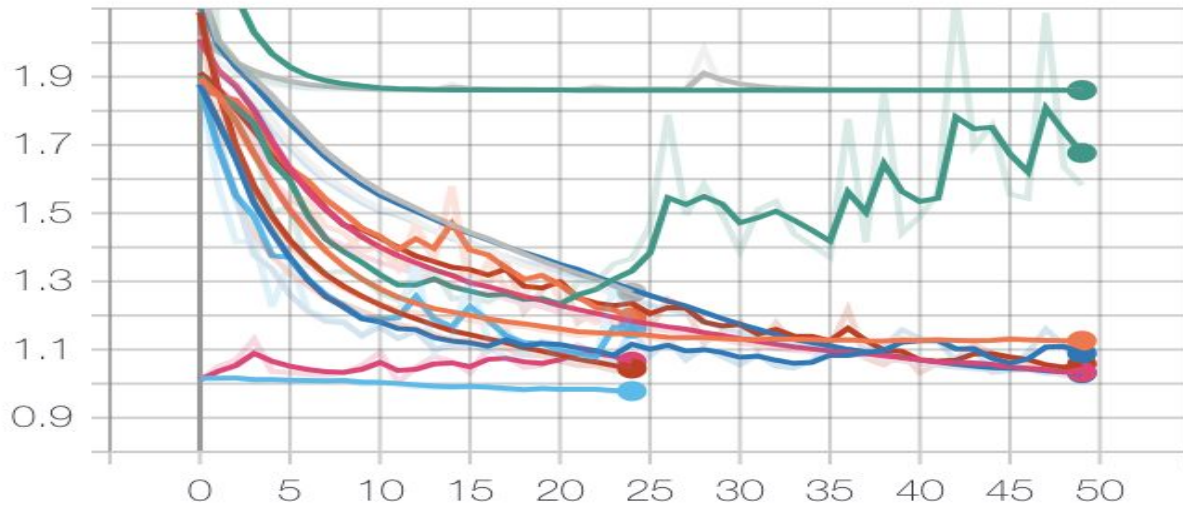


Tensorboard:

<https://tensorboard.dev/experiment/GkvcQqC8SI2UVUIHvxPaDg/#scalars>

Loss






epoch_loss



Tensorboard:

<https://tensorboard.dev/experiment/GkvcQqC8SI2UVUIHvxPaDg/#scalars>

Visualizing Edges



ResNet

Train on 29141 samples, validate on 3643 samples


```
Epoch 1/20
29141/29141 [=====] - 39s 1ms/step - loss: 0.4578 - accuracy: 0.8426 - val_loss: 1.9830 - val_accuracy: 0.5202
Epoch 2/20
29141/29141 [=====] - 38s 1ms/step - loss: 0.1191 - accuracy: 0.9632 - val_loss: 2.1124 - val_accuracy: 0.5597
Epoch 3/20
29141/29141 [=====] - 38s 1ms/step - loss: 0.0666 - accuracy: 0.9846 - val_loss: 2.2761 - val_accuracy: 0.5380
Epoch 4/20
29141/29141 [=====] - 38s 1ms/step - loss: 0.0552 - accuracy: 0.9895 - val_loss: 2.3251 - val_accuracy: 0.5284
Epoch 5/20
29141/29141 [=====] - 38s 1ms/step - loss: 0.0672 - accuracy: 0.9856 - val_loss: 2.6464 - val_accuracy: 0.4774
Epoch 6/20
29141/29141 [=====] - 38s 1ms/step - loss: 0.1354 - accuracy: 0.9559 - val_loss: 2.6523 - val_accuracy: 0.4996
Epoch 7/20
29141/29141 [=====] - 38s 1ms/step - loss: 0.1015 - accuracy: 0.9700 - val_loss: 2.4334 - val_accuracy: 0.5383
Epoch 8/20
```

OverFitting!!




Generative Adversarial Networks (GANs)

- Generates new images that look very realistic
- Two parts to GANs:
 - Generative network - generates candidates
 - Discriminative network - evaluates to real/fake
- Using a generative and discriminative network to generate faces would take a long time to generate several faces - [Colab](#)
- Used pretrained model and Nvidia's StyleGAN2 to generate very realistic faces - [Colab](#)



CNN Model on Faces using GANs



```
[ ] pred_df['label'].value_counts()
```

```
Happy      290  
Neutral    209  
Angry       1  
Name: label, dtype: int64
```

Hosting on TFX Serving Model

```
[ ] 1 %$bash --bg
    2 nohup tensorflow_model_server \
    3   --rest_api_port=8502 \
    4   --model_name=cnn_model_2 \
    5   --model_base_path="${MODEL_DIR}" >server.log 2>&1
```

☐ Starting job # 2 in a separate thread.

Make Predictions using TFX Serving Model

```
[ ] 1 import requests
2 headers = {"content-type": "application/json"}
3 json_response = requests.post('http://localhost:8502/v1/models/cnn_model_2:predict', data=data, headers=headers)
4 predictions = json.loads(json_response.text)['predictions']
5 output = predictions[0]
```

```
[ ] 1 from keras.preprocessing import image
2
3 classes = ['angry', 'disgust', 'fear', 'happy', 'sad', 'surprise', 'neutral', 'embarrassed']
4
5 img = image.load_img("/content/drive/My Drive/CMPE 258 - Deep Learning/258_final_project/data/stylegan2/image900.png", target_size=(48,48))
6 img = np.asarray(img)
7 plt.imshow(img)
8 img = np.expand_dims(img, axis=0)
9
10 print(output)
11 classes[np.argmax(output)]
```

```
☐ [0.0458349213, 3.94001631e-09, 0.012209137, 0.00210829754, 0.109127142, 2.6029853e-05, 0.830694497, 1.98055497e-14]
'neutral'
```



Conclusion

- A lot to learn about a human's face
- CNN was able to detect facial expression
 - Best accuracy: 64.53%
- Utilized GANs to help with testing
- TFX serving model to host our best model



References

[1]<https://arxiv.org/pdf/1902.01019.pdf>

[2]https://www.researchgate.net/publication/261498182_Face_expression_recognition_A_brief_overview_of_the_last_decade

[3]<https://www-sciencedirect-com.libaccess.sjlibrary.org/science/article/abs/pii/S0030402616302807>

[4]<https://www-sciencedirect-com.libaccess.sjlibrary.org/science/article/pii/S0925231217313644>

Dataset:<https://www.kaggle.com/c/challenges-in-representation-learning-facial-expression-recognition-challenge/data>





Thank you